

Assignment Four – Dynamic & Greedy

Ryan Munger
Ryan.Munger1@marist.edu

November 22, 2024

1 Introduction

1.1 Goals

Assignment 4 focuses on dynamic programming and greedy algorithms. First, I have to load in several weighted, directed graphs by parsing a file of instructions. Then, I must implement the Bellman-Ford dynamic programming algorithm for Single Source Shortest Path (SSSP) in order to find the optimal path from a source vertex to each other vertex in the graph. Graphs aside, the greedy algorithm I need to implement is a version of the fractional knapsack problem. I need to read in a file that contains information about spices (price, quantity, etc.) and I must conduct a spice heist on Arrakis for each knapsack provided, maximizing my take.

1.2 Write-up Format

In this report I will describe the logic being presented and the asymptotic running time of the algorithms implemented. Below the text explanation, relevant code will follow in C++.

1.3 Limerick of Luck

To maximally fill your knapsack
And remain unburdened with a fallback,
One must employ an algo of greed,
Your capacity will not exceed,
A heist for the ages, engraved on a plaque.

2 Weighted Directed Graphs

2.1 Reading the Instructions

I was provided a file of instructions to create weighted directed graphs such as *add vertex* and *add edge*. I was able to utilize most of the code from my previous assignment (with different regex) to create my graph. This time, I only had to create a linked object representation.

```
154 void createGraphs(const string& filename) {
155     int graphCount = 1;
156     cout << "Graph #" << graphCount << ":\n" << endl;
157     Graph currentGraph(to_string(graphCount));
158     regex newGraphRe(R"(new graph)");
159     regex addVertexRe(R"(add\s*vertex\s*(\S+))");
160     regex addEdgeRe(R"(add\s*edge\s*(\S+)\s*-\s*(\S+)\s*(-?\d+)");
161
162     ifstream file(filename); // input file stream
163     if (!file) {
164         cerr << "File opening failed." << endl;
165     }
166     string instruction;
167
168     while (getline(file, instruction)) {
169         // ignore any commands we don't know, empty lines, comments
170         // etc.
171         // regex will allow some slack with white space, but
172         // assuming perfect syntax by user
173         // case 1: start a new graph
174         if (regex_match(instruction, newGraphRe)) {
175             // check to see if the current graph has anything in it
176             // if not, no need to start a new one
177             if (!currentGraph.isEmpty()) {
178                 currentGraph.displayGraph();
179                 graphCount++;
180                 cout << "\n\nGraph #" << graphCount << ":\n" <<
181                 endl;
182                 currentGraph = Graph(to_string(graphCount)); //
183                 start a new graph
184             }
185             } else {
186                 smatch match; // captures subexpressions/groups
187
188                 if (regex_match(instruction, match, addVertexRe)) { //
189                     case 2: new vertex
190                     string newVertex = match[1].str();
191                     currentGraph.addVertex(newVertex);
192                 } else if (regex_match(instruction, match, addEdgeRe))
193                 { // case 3: new edge
194                     string v1 = match[1].str();
195                     string v2 = match[2].str();
196                     int weight = stoi(match[3].str());
197                     currentGraph.addEdge(v1, v2, weight);
198                 }
199             };
200     }
```

```

195     file.close();
196     currentGraph.displayGraph();
197 };

```

2.2 Graph Object

The graph object is even simpler than last time. Each graph has an ID and a map of linked vertex objects. This time, since my graph is weighted, I stored each edge in the neighbor list as a tuple. Adding vertices and edges is much simpler when we only have one representation to update! Since this is a directed graph, we only need to update one neighbor list. I made a new graph display function to ensure the graphs were created correctly.

```

19 struct linkedVertex {
20     string id;
21     int distance; // for SSSP
22     linkedVertex* predecessor; // for SSSP
23     vector<tuple<linkedVertex*, int>> neighbors; // no limit to
        neighbors!
24 };
25
26 void printLinkedVertex(linkedVertex v) {
27     cout << "LinkedVertex " << v.id << "; Neighbors: " << endl;
28     if (v.neighbors.empty()) {
29         cout << "\tNo Neighbors" << endl;
30     } else {
31         for (const auto& tuple : v.neighbors) {
32             cout << "\tVertex: " << get<vertexTupleIdx>(tuple)->id
33             << " Weight: " << get<weightTupleIdx>(tuple) << endl;
34         }
35     };
36
61     public:
62     Graph(string id) {
63         this->graphID = id;
64     };
65
66     void addVertex(string vertex) {
67         linkedObjs[vertex] = linkedVertex{vertex}; // store by
        value
68     };
69
70     void addEdge(string vertex1, string vertex2, int weight) {
71         this->linkedObjs[vertex1].neighbors.push_back(
        make_tuple(&linkedObjs[vertex2], weight));
72     };
73
74     bool isEmpty() {
75         return this->linkedObjs.empty();
76     };
77
78     void displayGraph() {
79         // print graph objects to ensure validity
80         for (const auto& pair : this->linkedObjs) {

```

```

81         printLinkedVertex(pair.second);
82     }
83
84     this->SSSP();
85 };

```

2.3 SSSP

The Single Source Shortest Path (SSSP) algorithm aims to find the shortest paths from a single source vertex to all other vertices in a graph. This is a powerful algorithm in scenarios like routing and navigation systems. Two important algorithms for SSSP are Dijkstra's algorithm and the Bellman-Ford algorithm, which I have implemented in this lab.

1. Initialize single source: $O(|V|)$.
2. Relax all edges $|V| - 1$ times: $O(|V| * |E|)$.
3. Check for negative weight cycles: $O(|E|)$.
4. Report the optimized paths (Not part of the algorithm directly).

I will explain the time complexities for each portion later. Combining the complexities of the algorithm's subroutines, the time complexity of Bellman-Ford is $O(|V| * |E|)$ where V is the set of vertices and E is the set of edges.

Initialize Single Source

1. Assign an initial distance of infinity to all vertices.
2. Set the source vertex distance to 0.

Why? This ensures that the shortest distance to the source vertex itself is 0, and all other vertices start with an "infinite" distance so that any path we compute is "cheaper." Since we traverse all the vertices once, this operation is $O(|V|)$.

```

42     void initSingleSource(linkedVertex* s) {
43         // set all vertices to distance infinite (large but not
max)
44         // no predecessors yet
45         for (auto& pair : this->linkedObjs) {
46             pair.second.distance = functionalInfinity;
47             pair.second.predecessor = nullptr;
48         }
49         // set single source
50         s->distance = 0;
51     };

```

Relaxing Edges

For each edge (*source*, *destination*) in the graph, check if the path from *source* to *destination* through that edge is better than the current distance to *destination*. If so, update *destination.distance* to the shorter value and set *source* as the predecessor of *destination*.

This process is called "relaxing" an edge. This is the main driver of Bellman-Ford because it allows us to find shorter paths and record them. We relax every edge $|V| - 1$ times because in the worst case, the optimal path can have $|V| - 1$ edges in it. Since we relax each edge essentially $|V|$ times, this is a costly operation at $O(|V| * |E|)$.

```
53         // find shortest path by recording the optimal choice
54         void relax(linkedVertex* source, linkedVertex* destination,
55             int weight) {
56             if (destination->distance > (source->distance + weight)
57         ) {
58                 destination->distance = (source->distance + weight)
59         ;
60                 destination->predecessor = source;
61             }
62         };
```

Why did the dynamic algorithm need scrap paper for its exam? *To write down all of its intermediate answers...*

Detecting Negative Weight Cycles

After we relax all of the edges $|V| - 1$ times, we must iterate over them once more. If we can still relax an edge, this indicates the presence of a negative weight cycle. We return false if this is the case, as it makes it impossible to reliably report shortest paths with this algorithm. A negative weight cycle occurs when there is a closed path (a loop or cycle) between two edges that has a negative cost. In this case, traveling around this loop over and over would reduce your cost indefinitely. The shortest path would be to follow this cycle infinite times before continuing on to your destination. Since this is a single traversal of the edges, this costs $O(|E|)$.

```
102         // detect negative weight cycles
103         for (auto& pair : this->linkedObjs) {
104             linkedVertex* current = &pair.second;
105             for (auto& edge : current->neighbors) {
106                 linkedVertex* destination = get<vertexTupleIdx>
107                 >(edge);
108                 int weight = get<weightTupleIdx>(edge);
109                 if (destination->distance > (current->distance
110                 + weight)) {
111                     return false; // negative weight cycle
112                     found
113                 }
114             }
115         }
```

```

90         // relax all edges |V| - 1 times
91         for (size_t i = 1; i < this->linkedObjs.size(); ++i) {
92             for (auto& pair : this->linkedObjs) {
93                 linkedVertex* current = &pair.second;
94                 for (auto& edge : current->neighbors) {
95                     linkedVertex* destination = get<
vertexTupleIdx>(edge);
96                     int weight = get<weightTupleIdx>(edge);
97                     relax(current, destination, weight);
98                 }
99             }
100         }

```

Report the Paths

Now that we have computed the shortest paths, we have to extract this data from our linked objects. The way to do this is simple: we can just check the predecessor of the destination, its predecessor, and so on until we are back at the source. Since this will be revealed in reverse order, pushing them onto a stack and then popping it will make them human readable. I did not count this as part of the algorithm's time complexity, as the necessary operations do not include this. However, this operation will take $O(|V|)$ since in the worst case we will have to pass through every vertex as a predecessor.

```

118         string getShortestPath(linkedVertex* destination) {
119             // follow predecessors (reverse order)
120             stack<string> pathStack;
121             linkedVertex* predecessor = destination;
122             while (predecessor != nullptr) {
123                 pathStack.push(predecessor->id);
124                 predecessor = predecessor->predecessor;
125             }
126
127             // put them in forward order for display
128             ostream pathstr;
129             while (!pathStack.empty()) {
130                 // Check if something is already in the stream for
->
131                 if (pathstr.tellp() > 0) {
132                     pathstr << "->";
133                 }
134                 pathstr << pathStack.top();
135                 pathStack.pop();
136             }
137             return pathstr.str();
138         };
139
140         void SSSP() {
141             linkedVertex* startVertex = &this->linkedObjs.begin()->
second;
142             // set distances, predecessors, etc
143             bellmanFord(startVertex);
144
145             cout << "SSSP: " << endl;
146             for (auto& pair : this->linkedObjs) {
147                 linkedVertex* current = &pair.second;

```

```

148         cout << startVertex->id << "->" << current->id << "
cost is " << setw(2) <<
149         current->distance << "; shortest path is " <<
getShortestPath(current) << endl;
150     }
151 };

```

Why did they add a timer to chess? *Mr. Dy Namic Algorithm...*

2.4 Graphs in Action

I have included one such graph below. This graph contains no negative weight cycles. As we can see, the graph was loaded in correctly and the shortest paths from the source to each other vertex was computed effectively.

```

3 Graph #1:
4
5 LinkedVertex 1; Neighbors:
6   Vertex: 2 Weight: 6
7   Vertex: 4 Weight: 7
8 LinkedVertex 2; Neighbors:
9   Vertex: 3 Weight: 5
10  Vertex: 4 Weight: 8
11  Vertex: 5 Weight: -4
12 LinkedVertex 3; Neighbors:
13  Vertex: 2 Weight: -2
14 LinkedVertex 4; Neighbors:
15  Vertex: 3 Weight: -3
16  Vertex: 5 Weight: 9
17 LinkedVertex 5; Neighbors:
18  Vertex: 3 Weight: 7
19  Vertex: 1 Weight: 2
20
21 SSSP:
22 1->1 cost is 0; shortest path is 1
23 1->2 cost is 2; shortest path is 1->4->3->2
24 1->3 cost is 4; shortest path is 1->4->3
25 1->4 cost is 7; shortest path is 1->4
26 1->5 cost is -2; shortest path is 1->4->3->2->5

```

3 Greedy Knapsack

Why wouldn't the greedy algorithm move?- *Staying local is important to him...*

3.1 Gathering Information

I gathered information about spices and knapsacks with regex in a similar fashion to my graphs. I stored my spices in a vector of Spice objects and my knapsacks in a vector as well. I used float values for everything, as this is the *fractional* knapsack problem, and there is no reason quantities and capacities cannot be decimal.

```
199 struct Spice {
200     string color;
201     float total_price;
202     float quantity;
203     float unit_price;
204 };

264 cout << "\n\nLoading in Spices and Knapsacks!" << endl;
265 regex spiceRe(R"(\s*spice\s*name\s*=\s*(\S*)\s*;\s*total_price\s*=\s*(\d*.\d*)\s*;\s*qty\s*=\s*(\d*.\d*)\s*;)" );
266 regex knapsackRe(R"(knapsack\s*capacity\s*=\s*(\d*.\d*)\s*;)" );
267
268 // store spices and knapsacks
269 vector<Spice> spiceInventory;
270 vector<float> knapsacks;
271
272 ifstream file(filename); // input file stream
273 if (!file) {
274     cerr << "File opening failed." << endl;
275 }
276 string instruction;
277
278 while (getline(file, instruction)) {
279     smatch match; // captures subexpressions/groups
280     // case 1: adding a spice
281     if (regex_match(instruction, match, spiceRe)) {
282         string color = match[1].str();
283         float total_price = stof(match[2].str());
284         float quantity = stof(match[3].str());
285         float unit_price = total_price / quantity;
286
287         Spice newSpice = Spice{color, total_price, quantity,
288                                unit_price};
289         printSpice(newSpice);
290         spiceInventory.push_back(newSpice);
291     } else if (regex_match(instruction, match, knapsackRe)) {
292         // case 2: knapsack
293         float newKnapsackCapacity = stof(match[1].str());
294         knapsacks.push_back(newKnapsackCapacity);
295     }
296 }
```



```

293         cout << "New Knapsack: " << newKnapsackCapacity << endl
294     };
295 };
296 file.close();
297 // sort our spices based on unit price
298 spiceSort(spiceInventory);
299 // maximize take for each knapsack!
300 cout << "\nMaximizing Take:" << endl;
301 for (float knapsack : knapsacks) {
302     maximizeTake(knapsack, spiceInventory);
303 }
304 };

```

3.2 Organizing Spice

To maximize take, we will examine the unit price of each spice (how much it is worth per quantity). To do this, we first sort the Spice list. I made a custom version of insertion sort (for simplicity) to accomplish this. I put it in descending order. Since I used insertion sort, this action will take $O(n^2)$ time due to the nested loop. We can use a better sorting algorithm, such as merge or quick sort, to optimize this down to $O(\log(n))$.

```

214 // Insertion sort to get descending order based on unit price
215 void spiceSort(vector<Spice>& arr) {
216     int n = arr.size();
217     for (int i = 1; i < n; i++) {
218         int insertIdx = i;
219         Spice currentCheck = arr[i];
220         for (int j = i-1; j >= 0; j--) {
221             if (arr[j].unit_price < currentCheck.unit_price) {
222                 arr[j+1] = arr[j];
223                 insertIdx = j;
224             } else {
225                 break;
226             }
227         }
228         arr[insertIdx] = currentCheck;
229     }
230 };

```

3.3 Maximizing Take

I implemented a greedy algorithm. This class of algorithm takes locally optimal choices and hopes for a globally optimal solution. In this case, we will achieve a globally optimal solution by pillaging as much of the highest value spice we can fit, then the next, and so on. This algorithm will only take $O(n)$ time as it is simply a single traversal of the spice list. It is even less than a single traversal, as we can expect most knapsacks to fill up before we reach the end of our spice inventory list! Thus, fractional knapsack is a $O(n \log(n))$ algorithm if you count sorting the spice list, or $O(n)$ on its own.

1. Examine the most valuable spice.
2. If we have no more knapsack capacity, we are done.
3. If we have more capacity than quantity of that spice, take everything! Record our scoops.
4. If we have less capacity than the quantity of that spice, take as much as we can fit. Record scoops.
5. Move to the next most valuable spice and repeat.
6. Finally, report on our knapsack value and scoops taken.

```
232 void maximizeTake(float knapsack, vector<Spice> spices) {
233     float knapValue = 0;
234     if (knapsack == 0) {
235         cout << "Knapsack of Capacity " << fixed << setprecision(2)
236             << knapsack << " is worth " <<
237             fixed << setprecision(2) << knapValue << " quatloos and
238             contains no scoops." << endl;
239         return;
240     }
241     ostringstream scoops;
242     float capacityLeft = knapsack;
243     for (Spice spice : spices) {
244         if (capacityLeft == 0) {
245             break; // no more spice!!
246         } else if (capacityLeft >= spice.quantity) { // take all the
247             spice
248                 capacityLeft -= spice.quantity;
249                 knapValue += spice.total_price;
250                 scoops << fixed << setprecision(2) << spice.quantity <<
251                 " scoops of " << spice.color << ", ";
252             } else if (capacityLeft < spice.quantity) { // take what we
253                 can fit
254                     knapValue += capacityLeft * spice.unit_price;
255                     scoops << fixed << setprecision(2) << capacityLeft << "
256                     scoops of " << spice.color << ", ";
257                     capacityLeft = 0;
258                 }
259     }
260     string scoopString = scoops.str();
```

```

255 // replace last comma with period
256 scoopString.pop_back();
257 scoopString.back() = '.';
258
259 cout << "Knapsack of Capacity " << fixed << setprecision(2) <<
knapsack << " is worth " <<
260 fixed << setprecision(2) << knapValue << " quatloos and
contains " << scoopString << endl;
261 };

```

3.4 Greed in Action

Spices were loaded, knapsacks were created, and heists were completed! For each knapsack, the most optimal scoops were scooped. I think "scamped" would be cooler to say though. Optimal scoops were scamped.

```

123 Loading in Spices and Knapsacks!
124 Spice:
125   Color: red
126   Total Price: 4
127   Quantity: 4
128   Unit Price: 1
129 Spice:
130   Color: green
131   Total Price: 12
132   Quantity: 6
133   Unit Price: 2
134 Spice:
135   Color: blue
136   Total Price: 40
137   Quantity: 8
138   Unit Price: 5
139 Spice:
140   Color: orange
141   Total Price: 18
142   Quantity: 2
143   Unit Price: 9
144 New Knapsack: 1
145 New Knapsack: 6
146 New Knapsack: 10
147 New Knapsack: 20
148 New Knapsack: 21
149
150 Maximizing Take:
151 Knapsack of Capacity 1.00 is worth 9.00 quatloos and contains 1.00
scoops of orange.
152 Knapsack of Capacity 6.00 is worth 38.00 quatloos and contains 2.00
scoops of orange, 4.00 scoops of blue.
153 Knapsack of Capacity 10.00 is worth 58.00 quatloos and contains
2.00 scoops of orange, 8.00 scoops of blue.
154 Knapsack of Capacity 20.00 is worth 74.00 quatloos and contains
2.00 scoops of orange, 8.00 scoops of blue, 6.00 scoops of
green, 4.00 scoops of red.
155 Knapsack of Capacity 21.00 is worth 74.00 quatloos and contains
2.00 scoops of orange, 8.00 scoops of blue, 6.00 scoops of
green, 4.00 scoops of red.

```

4 Conclusion

Dynamic programming is extremely powerful, yet not very efficient. Dynamic algorithms such as Bellman-Ford for SSSP tackle variable & complex problems with ease! It is interesting to think about algorithms such as these that run our navigation systems, networking, and more!

Greedy algorithms are much simpler than they seem. All that needs to be done is to take the greediest, most locally & immediately optimal action. It is important to note though that while this did produce a globally optimal solution in this knapsack case, it does not always turn out this way, such as the cases of the 0-1 knapsack problem and traversing directed graphs!

Why did the greedy algorithm get full so quickly? *It ate all the appetizers and spared no room!*
